Pranav Mishra 773-280-4615
pmishr23@uic.edu github.com/PranavMishra17
linkedin.com/in/pranavgamedev portfolio-pranav-mishra-paranoid.vercel.app

Dear Hiring Team,

I build AI systems for a living, and I want to do that work at the frontier. I am writing because the work you do sits exactly where I aim to spend my career: production agentic systems that real users depend on, not demos.

I am currently a Founding LLM Engineer at alfred_ (Techstars-backed, NYC), where I design the evaluation harness for a multi-agent SMS assistant serving 5K+ active users – trace replay, verdict-regression detection, and tool-calling reliability across a TypeScript backend. Before that, I shipped MockFlow-AI, a real-time voice interview platform with FSM-driven multi-agent orchestration, streaming STT/TTS, and sub-400ms end-to-end latency (mockflow-ai.onrender.com). My research on MetaRAG (accepted at IEEE CAI 2026) hit 82.5% precision and 0.925 Hit@10 by combining LLM-generated metadata, cross-encoder reranking, and TF-IDF embeddings – deployed on AWS SageMaker at p99 under 300ms.

What pulls me to your team is the chance to push agentic AI past plumbing into reliability. I built SoulEngine, a multi-provider LLM agent framework with MCP tool invocation and dual-instance memory architecture, because I wanted to understand where current agent designs break. My published work on TeamMedAgents and SLM-TeamMedAgents (ICML 2026 submission) explored multi-agent teamwork components – consensus voting, role specialization, reasoning traces – and hit 77.63% accuracy with a 3.1x inference speedup over frontier LLMs.

I am a builder. I ship vertical slices: research to runtime, prototype to production. I do not wait for permission to figure something out. If your team is solving hard problems in agentic systems, LLM reliability, or evaluation infrastructure, I would love to be part of it.

Thanks for your time, and I look forward to talking.

Sincerely,
Pranav Mishra