

Pranav Mishra
pmishr23@uic.edu
linkedin.com/in/pranavgamedev

773-280-4615
github.com/PranavMishra17
portfolio-pranav-mishra-paranoid.vercel.app

Dear Hiring Team,

I am an ML engineer who likes the full arc – from training loop to inference latency budget. I am writing because I want to spend the next chapter of my career at the frontier of applied ML, on systems that actually ship and stay shipped.

My published research on MetaRAG (accepted at IEEE CAI 2026) achieved 82.5% precision and 0.925 Hit@10 by combining LLM-generated metadata, cross-encoder reranking, and TF-IDF-weighted embeddings; I reduced hallucinations 25% via an automated evaluation framework, and deployed the production system on AWS SageMaker with MLflow tracking, handling 10k queries/day at p99 under 300ms. My TeamMedAgents paper (PAKDD 2026, with SLM-TeamMedAgents at ICML 2026 submission) optimized 4B parameter SLMs in a multi-agent reasoning system – 77.63% accuracy across 8 medical benchmarks with a 3.1x inference speedup over frontier LLMs through trust-weighted consensus voting and reasoning traces. I presented MetaRAG to 700+ analytics professionals at INFORMS Analytics+.

On the systems side, I built SnakeAI-MLOps, a reinforcement-learning framework in C++ with LibTorch – reward shaping, experience replay, hyperparameter optimization, TorchScript inference at $\sim 140 \mu s$, and a 50,000-iteration soak validation – wrapped in an MLOps pipeline with Docker CI/CD, model versioning, and deployment monitoring. As Founding LLM Engineer at alfred_, I am designing the production eval harness for an SMS agent serving 5K+ users: trace replay, fixture-driven scenarios, and tool-calling regression detection.

I am a builder. I write the training code, I write the inference code, I write the eval harness, and I care about the seam between them. If your team is solving hard problems in applied ML – production inference, eval infrastructure, retrieval, multi-agent reasoning – I would love to contribute.

Thanks for your time, and I look forward to talking.

Sincerely,
Pranav Mishra